

CS 580 RYL 3

```
In [ ]: # dependencies
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
# import sympy as sp
import scipy
# from PIL import Image


# Zac noted in recitation that we could use official package documentation
# on the RYL's moving forward'. I did that for this assignment.
```

```
In [ ]: # setup
# The file babyboomerdivorce.csv contains
# data on the rate of baby-boomer divorce.
# The independent variables are marriage age (x) and years of school (y).

df = pd.read_csv("./babyboomerdivorce.csv")
x = df.age_at_marriage.values
y = df.years_of_education.values
z = df.divorce_rate_per_100.values
```

```
In [3]: # visualize
plt.scatter(x, z, color="r")
plt.scatter(y, z, color="b")
```

```
Out[3]: <matplotlib.collections.PathCollection at 0x1140265d0>
```



```
In [4]: # A
# (a) Formulate regression problems for the following models:
#  $z(x,y) = \theta_{00} + \theta_{10}x + \theta_{01}y$ 
#  $z(x,y) = \theta_{00} + \theta_{10}x + \theta_{01}y + \theta_{11}xy$ 
#  $z(x,y) = \theta_{00} + \theta_{10}x + \theta_{01}y + \theta_{11}xy + \theta_{20}x^2 + \theta_{02}y^2$ 

ones = np.ones(df.shape[0])
A1 = np.array([ones, x, y]).T
A2 = np.array([ones, x, y, x * y]).T
A3 = np.array([ones, x, y, x * y, x**2, y**2]).T
```

```
In [5]: # B
# (b) Solve each regression and report best-fit parameters and uncertainties.
theta_hat_1, sse1, rank, s = np.linalg.lstsq(A1, z)
theta_hat_2, sse2, rank, s = np.linalg.lstsq(A2, z)
theta_hat_3, sse3, rank, s = np.linalg.lstsq(A3, z)
```

```

print("-----")
print(f"the model for z(x,y) = 000 + 010x+ 001y gives theta_hat:\n{theta_hat_1}")
print(f"SSE is {sse1}")
print("-----")
print(f"the model for z(x,y) = 000 + 010x+ 001y+ 011xy gives theta_hat:\n{theta_hat_2}")
print(f"SSE is {sse2}")
print("-----")
print(f"the model for 000 + 010x+ 001y+ 011xy+ 020x2 + 002y2 gives theta_hat:\n{theta_hat_3}")
print(f"SSE is {sse3}")
print("-----")

```

```

-----
the model for z(x,y) = 000 + 010x+ 001y gives theta_hat:
[122.18308858 -1.88084719 -1.94914286]
SSE is [387.28474669]

```

```

-----
the model for z(x,y) = 000 + 010x+ 001y+ 011xy gives theta_hat:
[ 1.67870700e+02 -3.34051530e+00 -5.10001262e+00  1.00666766e-01]
SSE is [250.80275913]

```

```

-----
the model for 000 + 010x+ 001y+ 011xy+ 020x2 + 002y2 gives theta_hat:
[ 1.14478739e+02 -2.17977168e+00 -3.05970381e-02  1.00666766e-01
 -1.87273229e-02 -1.67386364e-01]
SSE is [191.97966108]
-----

```

```

In [6]: # C
# (c) Compute leave-one-out cross validation error for each model and compare.
name1 = "z(x,y) = 000 + 010x+ 001y" # cross validation error: 2698.0644755999683
name2 = "z(x,y) = 000 + 010x+ 001y+ 011xy" # cross validation error: 1647.6357933662268
name3 = "000 + 010x+ 001y+ 011xy+ 020x2 + 002y2" # cross validation error: 1120.1853500339723

for nm, AA in [(name1, A1), (name2, A2), (name3, A3)]:
    print("-----")
    print(f"leave-one-out cross validation for {nm}")
    errs = []
    for i in range(AA.shape[0]):
        tmp_a = np.concatenate([AA[:i], AA[i+1:]])
        left_out_a = AA[i]
        tmp_z = np.concatenate([z[:i], z[i+1:]])
        left_out_z = z[i]

```

```

    theta_hat_tmp, _, _, _ = np.linalg.lstsq(AA, z)
    err = np.linalg.norm((tmp_z - np.matmul(tmp_a, theta_hat_tmp))**2, ord=2)
    errs.append(err)
tot_err = sum(errs)
print(f"cross validation error: {tot_err}")
print("-----")

```

Comparison:

Here we find that the larger-dimensional bases give improved error in cross-validation.

This suggests that these models with more parameters generalize better.

```

-----
leave-one-out cross validation for  $z(x,y) = \theta_{00} + \theta_{10}x + \theta_{01}y$ 
cross validation error: 2698.0644755999683
-----

```

```

-----
leave-one-out cross validation for  $z(x,y) = \theta_{00} + \theta_{10}x + \theta_{01}y + \theta_{11}xy$ 
cross validation error: 1647.6357933662268
-----

```

```

-----
leave-one-out cross validation for  $\theta_{00} + \theta_{10}x + \theta_{01}y + \theta_{11}xy + \theta_{20}x^2 + \theta_{02}y^2$ 
cross validation error: 1120.1853500339723
-----

```

```

In [ ]: # D
# (d) Solve the model  $z(x,y) = \theta_{00} + \theta_{10}x + \theta_{01}y$  using the 1-norm. Report parameters.

# up to now we have been solving for theta with least squares:
theta_hat_eg = ((A1.T@A1)**-1)@A1.T@z
# because the above is not numerically stable, we solve with numpy:
theta_hat_eg, _, _, _ = np.linalg.lstsq(A1, z)

# however, the expression `((A1.T@A1)**-1)@A1.T@z` is derived using the L-2 norm.
# If we want a solution based on a different norm, we need a different strategy
#  $z(x,y) = \theta_{00} + \theta_{10}x + \theta_{01}y$  is linear, so I think a closed-form expression could
# be derived using matrix calculus as with the least-squares solution
# I tried to work this out by hand, but couldn't find the solution.
# I can also find the value algorithmically as follows.
# this will not be absolutely precise, but will achieve a practical level of certainty.

def normed_err_from_thetas(thet):
    """ objective function to minimize """

```

```
z_est = np.matmul(A1, thet)
e_vec = z-z_est
normed_e = np.linalg.norm(e_vec, ord=1)
return normed_e

# use theta_hat_1 as the initial value for scipy's optimization algorithm
norm1_theta_hat_1 = scipy.optimize.minimize(normed_err_from_thetas, theta_hat_1).x

print(f"My solution is {norm1_theta_hat_1}")
```

My solution is [126.13757252 -1.77509354 -2.37689689]

580
9-27-25

ZyL3, D

$$\begin{bmatrix} z \\ z \\ z \\ z \end{bmatrix} = \begin{bmatrix} A_1 & A_2 & A_3 \\ A_1 & A_2 & A_3 \end{bmatrix} \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} e_1 \\ e_2 \\ e_3 \\ e_4 \end{bmatrix}$$

$$z(x, y) = \theta_0 + \theta_1 x + \theta_2 y + e$$

$$Z = A = \begin{bmatrix} 1 & x & y \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \end{bmatrix}$$

$$Z = A\theta + e \Rightarrow e = Z - A\theta$$

Objective: $\hat{\theta} = \underset{\theta}{\text{argmin}} \|Z - A\theta\|_1$

P-Norm: $\left(\sum_i |e_i| \right)^{1/L}$

LSTD: $\sqrt{x^2 + y^2}$

$$P_1 = \sum_i e_i = \sum_i (z_i - A_i \theta) = \left(\sum_i z_i - \sum_j (A_{ij}) (\theta_j) \right)$$

WE WANT TO MINIMIZE THIS

$$P = \sum_i (z_i - A_{i0}\theta_0 - A_{i1}\theta_1 - A_{i2}\theta_2)$$

$$\frac{\partial P}{\partial \theta_0} = \sum_i -A_{i0}$$

$$\frac{\partial P}{\partial \theta_1} = \sum_i -A_{i1}$$

$$\frac{\partial P}{\partial \theta_2} = \sum_i -A_{i2}$$

I KNOW THIS
ISN'T RIGHT.
I'LL FORGIVE INSTAD.

580
9-27-25

2422 NOTES

$$LSQS = (A'A)^{-1} A' Y$$

$$\hat{\theta} = \underset{\theta}{\text{Argmin}} \|e\|_2 = \underset{\theta}{\text{Argmin}} \|Y - A\theta\|_2$$

$$H \text{ ASSUMES } \|e\|_2^2 : \frac{1}{2} A'A, \frac{\partial^2 L}{\partial \theta_i \partial \theta_k}$$

$$\langle U, V \rangle = U'GV \quad \text{RANK: LINEARLY INDEP.}$$

SINGULAR: SQUARE BUT NOT INVERTIBLE

THEOREM INV. MAT.

PROB.

$$P(B|A) = \frac{P(A \cap B)}{P(A)}$$

BAYES.

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

TOTAL PROB.

$$P(A) = \sum P(A|B_n)P(B_n)$$

$n \times n$, IFF:

- A^{-1} EXISTS
- $\det(A) \neq 0$
- $\text{RANK} = n$
- ROWS + COLS LIN IND.
- $Ax = 0$ ONLY TRIVIAL
- $Ax = b$ UNIQUE
- $A \rightarrow REF \rightarrow I$
- $A \rightarrow REF \rightarrow \text{PIVOTS}$
- $0 \notin \{\lambda_{1,2}\}$

$$\text{LIKELIHOOD: } P(\vec{y}|\theta)$$

PROBABILITY: ROW

MLE: FIND LOSS, $\frac{dL}{d\lambda}$, SET TO 0, SOLVE FOR λ , $\hat{\lambda}_{MLE}$ (?)