

# CS 580 RYL 4

## Setup

```
In [1]: # dependencies
# import warnings

import jax
import jax.numpy as jnp
# import numpy as np
# import pandas as pd
# import PIL
# import scipy
# import sympy as sp
from matplotlib import pyplot as plt

# warnings.filterwarnings("ignore")

# Zac noted in recitation that we could use official package documentation
# on the RYL's moving forward'. I did that for this assignment.
```

## RYL 4

```
In [2]: # Consider the model:
# y= 01 * e^(-02x) + 03
# and the data points
# {(0, 2.300), (1, 1.036), (2, 0.571)}

def model(x, th):
    """
    x: scalar or vector
    th: 3-dim vector
    """
    return (th[0] * jnp.exp(-th[1]*x)) + th[2]


data = jnp.array([
    [0, 2.300],
    [1, 1.036],
    [2, 0.571]
])

data_x = data[:, 0]
data_y = data[:, 1]
```

```
In [3]: # plot a generic version of our model with the data for context

dummy_x = jnp.linspace(-1,10,100)
dummy_y = model(dummy_x, jnp.array([1.5,0.6,0])) # using theta from part c
plt.plot(dummy_x, dummy_y)
plt.scatter(data_x, data_y, marker="x", color="black")
```

```
Out[3]: <matplotlib.collections.PathCollection at 0x1153c8ec0>
```



In [4]: *# (a) Define your cost as the square of the l2 norm and write a function which takes vectors x,y, # and  $\theta$  and returns the cost.*

```
def cost(x, y, th):
    """
    x: vector
    y: vector (same shape as x)
    th: 3-dim vector
    """
    return jnp.linalg.norm(y-model(x, th), ord=2)**2
```

In [5]: *# (b) Use an automatic differentiation library of your choice to take the gradient of your cost # function*

```
cost_gradient = jax.grad(cost, 2) # gradient with respect to theta
```

In [6]: *# (c) Calculate the normalized negative gradient ( $\hat{g}$ ) at  $\theta=(1.5, 0.6, 0)$*

```
theta_0 = jnp.array([1.5, 0.6, 0])
g_hat = jnp.linalg.norm(-cost_gradient(data_x, data_y, theta_0), ord=2)
g_hat
```

Out[6]: Array(3.0126598, dtype=float32)


In [7]: *# (d) Plot  $L(\theta_0 + \alpha \hat{g})$  where  $L$  is your loss function,  $\theta_0$  is the theta vector given in # part c,  $\hat{g}$  is the normalized negative gradient from part c, and  $\alpha$  is a scalar with # values ranging from 0 to 1.*

```
def L(alpha):
    return [cost(data_x, data_y, theta_0 + a * g_hat) for a in alpha]

alpha = jnp.linspace(0, 1, 100)
losses = L(alpha)

plt.plot(alpha, losses)
plt.ylim(-10,40)
```

Out[7]: (-10.0, 40.0)



In [8]: # (e) Using the same y-axis, plot  $L(\theta_0 + \alpha d_{\text{hat}})$  where  $d_{\text{hat}} = (0.138, 0.983, 0.118)$

```
def L_g(alpha):
    return [cost(data_x, data_y, theta_0 + a * g_hat) for a in alph]


d_hat = jnp.array([0.138, 0.983, 0.118])

def L_d(alpha):
    return [cost(data_x, data_y, theta_0 + a * d_hat) for a in alph]

alpha = jnp.linspace(0, 1, 100)
losses_g_hat = L_g(alpha)
losses_d_hat = L_d(alpha)

plt.plot(alpha, losses_g_hat, label="losses_g_hat")
plt.plot(alpha, losses_d_hat, label="losses_d_hat")
plt.ylim(-10, 40)
plt.legend()
```

Out[8]: <matplotlib.legend.Legend at 0x1153c96a0>



In [9]: # (f) Using your two plots, explain why the gradient is useful in the context of modeling.

```
# the gradient is useful in the context of modeling because stepping around our cost surface with  
# the gradient with respect to the parameters allows characterization of the cost surface, and  
# can therefore help us to find minima.
```